

# Naïve Bayesian Classifier (method review)

Grishin Andrey (group: M05-318a)

Department of Machine Learning and Digital Humanities

Moscow Institute of Physics and Technology

November 11, 2023

# Plan

- 1 Introduction
  - What is Bayesian Classifier?
  - Why (what for) do we need it?
- 2 Maths of Bayesian Classifier
  - Model assumptions
  - Main maths
- 3 Iris classification problem
  - Problem state
  - Fitting results
- 4 Summary
  - References
  - Conclusion

# What is Bayesian Classifier?

Bayesian Classifier — **probabilistic** classification model:

- **Foundation:** Bayes's theorem.
- **Authors:** Thomas Bayes, Richard Price.
- **Title:** “An Essay towards Solving a Problem in the Doctrine of Chances”<sup>1</sup>.
- **Initial use:** Estimate distribution parameters.
- **Year:** 1763

---

<sup>1</sup>Reference to the article [Bayes and Price, 1763].

# Why (what for) do we need it?

- **Classifier:**

- **Segregate** data observations.
- **Baseline** classification models.

- **Theorem:**

- **Estimate** distribution parameters (normal, binomial, etc.).
- **Estimate** posterior probabilities of events.

## Model assumptions

**Let**  $X \in \mathbb{R}^{n \times m}$ ,  $y \in \{C_1, C_2, \dots, C_k\}$ . **Then** need to find  $f(\cdot)$  such as:

$$\mathbb{E}(y | X) = f(X) \Leftrightarrow y = f(X) + \epsilon, \quad \epsilon \sim \mathcal{N}(0, \sigma^2) \quad (1)$$

$f(\cdot)$  = Bayesian Classifier.  $X_j \in \mathbb{R}^{1 \times m} : j = \overline{1, n}$  is r.v.  $\mu$  and  $\Sigma$  estimation – tough.

**Assume:**

- $X_j \sim \mathcal{N}(\mu, \Sigma) : \mu \in \mathbb{R}^{m \times 1}, \Sigma \in \mathbb{R}^{m \times m}$ .
- $X_j = [\xi_1, \dots, \xi_m] : \xi_i \stackrel{\text{i.i.d.}}{\sim} \mathcal{N}(\mu_i, \sigma_i^2) \Rightarrow \Pr[X_j] = \prod_{i=1}^m \Pr[\xi_i]$ .

**Note:** classification  $\approx$  regression: probability  $\in [0, 1]$ .

## Main maths I

**Bayesian theorem** is formalized like this:

$$\Pr[A | B] = \frac{\Pr[B | A] \times \Pr[A]}{\Pr[B]} \quad (2)$$

**Bayesian classifier**, where  $c_t : t = \overline{1, k}$  is the class number:

$$\Pr[y_j = c_t | X_j] = \frac{\Pr[X_j | y_j = c_t] \times \Pr[y_j = c_t]}{\Pr[X_j]} \quad (3)$$

The formula (3) can be respectively expressed as:

$$\text{posterior} = \frac{\text{likelihood} \times \text{prior}}{\text{evidence}} \quad (4)$$

## Main maths II

**But** evidence  $\Pr[X_j]$  is the same  $\forall c_t : t = \overline{1, k}$ . Hence, we do not strictly need it.

$$\Pr[y_j = c_t \mid X_j] \propto \Pr[X_j \mid y_j = c_t] \times \Pr[y_j = c_t] \quad (5)$$

**After expansion** of the (5) we almost have the final formula:

$$\Pr[y_j = c_t \mid X_j] \propto \left( \prod_{i=1}^m \Pr[X_{ji}] \right) \times \Pr[y_j = c_t] \quad (6)$$

**So**, (6) gives us the expression for  $X_j$ 's class:

$$\hat{C}_j = \arg \max_{t=\overline{1, k}} \left\{ \left( \prod_{i=1}^m \Pr[X_{ji}] \right) \times \Pr[y_j = c_t] \right\} \quad (7)$$

## Main maths III

**Q:** What exactly is  $\prod_{i=1}^m \Pr[X_{ji}]$  and how to estimate  $(\mu_j, \sigma_j^2)$ ?

**A:** To make things clear, let's reduce the number of **features** and **classes** to 1.

**Assumed**  $\mathcal{N}(\mu, \sigma^2) \Rightarrow 2$  parameters to be estimated. MLE method [Wilks, 1938].

$$L = \prod_{j=1}^n \Pr[X_j | \mu, \sigma] \equiv \prod_{j=1}^n L(\mu, \sigma | X_j) \quad (8)$$

No doubts to get the following optimization problem:

$$\hat{\mu}, \hat{\sigma} = \arg \max_{\mu, \sigma} \left\{ \prod_{j=1}^n L(\mu, \sigma | X_j) \right\} : \mu \in \mathbb{R}, \sigma \in \mathbb{R}_+ \quad (9)$$



## Main maths IV

**Apply**  $\ln(x)$  as extrema is invariant to monotonic increasing transforms.

$$\arg \max_{\mu, \sigma} \left\{ \prod_{j=1}^n L(\mu, \sigma | X_j) \right\} \sim \arg \max_{\mu, \sigma} \left\{ \ln \left( \prod_{j=1}^n L(\mu, \sigma | X_j) \right) \right\} \quad (10)$$

It leads to simplification:

$$\arg \max_{\mu, \sigma} \left\{ \sum_{j=1}^n \ln \left( L(\mu, \sigma | X_j) \right) \right\} \quad (11)$$

**Finally** to estimate  $\mu$  and  $\sigma$  we have to optimize  $L(\mu, \sigma) = \sum_{j=1}^n \ln \left( \Pr[X_j | \mu, \sigma] \right)$ .

## Main maths V

And the result is the following (**unbiased** variance):

$$\max_{\mu, \sigma} \left\{ \sum_{j=1}^n \ln \left( \frac{1}{\sqrt{2\pi\sigma^2}} \exp \left( \frac{(X_j - \mu)^2}{2\sigma^2} \right) \right) \right\} \Rightarrow \begin{cases} \hat{\mu} &= \frac{\sum_{j=1}^n X_j}{n} \\ \hat{\sigma}^2 &= \frac{\sum_{j=1}^n (X_j - \bar{X})^2}{n-1} \end{cases} \quad (12)$$

## Main maths VI

With  $\mathbf{m}$  i.i.d. features we obtain the following estimation ( $i = \overline{1, m}$ ):

### Mean estimation

$$\hat{\mu}_i = \frac{\sum_{j=1}^n X_{ji}}{n} \quad (13)$$

### Variance estimation

$$\hat{\sigma}_i^2 = \frac{\sum_{j=1}^n (X_{ji} - \bar{X}_i)^2}{n - 1} \quad (14)$$

## Main maths VII

With  $m$  features and  $k$  classes estimation (13) (14) turns into ( $i = \overline{1, m}, t = \overline{1, k}$ ):

### Mean estimation

$$\hat{\mu}_i(c_t) = \frac{\sum_{j=1}^n X_{ji} [y_j = c_t]}{\sum_{j=1}^n [y_j = c_t]} \quad (15)$$

### Variance estimation

$$\hat{\sigma}_i^2(c_t) = \frac{\sum_{j=1}^n (X_{ji} - \hat{\mu}_i(c_t))^2 [y_j = c_t]}{\left(\sum_{j=1}^n [y_j = c_t]\right) - 1} \quad (16)$$

## Main maths VIII

For **simplicity**, Gaussian NB works if  $(X_{ji}(c_t): \mathbf{j}$  obs.,  $\mathbf{i}$  feat.,  $\mathbf{c}_t$  class):

Naïve Bayesian classifier requirement

$$\forall \mathbf{c}_t, \mathbf{i} : \left\{ \begin{array}{l} t = \overline{1, k} \\ i = \overline{1, m} \end{array} \right. \exists \hat{\mu}_i(c_t) \wedge \hat{\sigma}_i^2(c_t) \hookrightarrow X_{ji}(c_t) \sim \mathcal{N}(\hat{\mu}_i(c_t), \hat{\sigma}_i^2(c_t)) \quad (17)$$

## Main maths IX

Multiple distributions can be used:

- Normal, Binomial.
- Laplacian, Exponential.
- Rayleigh's distribution.
- etc.

## Problem state I

Obtained  $X \in \mathbb{R}^{150 \times 4}$ ,  $y \in \{0, 1, 2\}$ . Classification problem: iris dataset.

Code	Name
0	setosa
1	versicolor
2	virginica

Table 1: Iris encoding

Code	Name	Measure
0	sepal length	cm
1	sepal width	cm
2	petal length	cm
3	petal width	cm

Table 2: Feature names

Class	# obs.
setosa	50 (0.3)
versicolor	50 (0.3)
virginica	50 (0.3)

Table 3: Class balancing

## Problem state II




Figure 1: Sepal length

Name	Value
mean	5.01
std	0.35
max	5.80
75%	5.20
50%	5.00
25%	4.80
min	4.30

Table 4: Stats.




Figure 2: Distribution



## Problem state III




Figure 3: Sepal width

Name	Value
mean	3.43
std	0.38
max	4.40
75%	3.68
50%	3.40
25%	3.20
min	2.30

Table 5: Stats.




Figure 4: Distribution

## Problem state IV




Figure 5: Petal length

Name	Value
mean	1.46
std	0.17
max	1.90
75%	1.58
50%	1.50
25%	1.40
min	1.00

Table 6: Stats.




Figure 6: Distribution

## Problem state V




Figure 7: Petal width

Name	Value
mean	0.25
std	0.11
max	0.60
75%	0.30
50%	0.20
25%	0.20
min	0.10

Table 7: Stats.

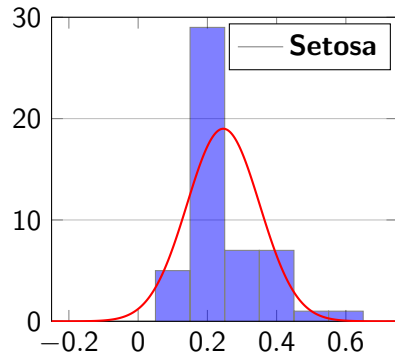


Figure 8: Distribution

## Problem state VI




Figure 9: Sepal length

Name	Value
mean	5.94
std	0.52
max	7.00
75%	6.30
50%	5.90
25%	5.60
min	4.90

Table 8: Stats.

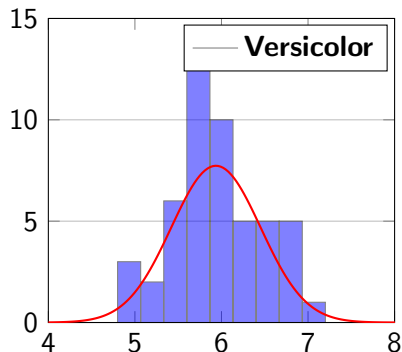


Figure 10: Distribution

## Problem state VII




Figure 11: Sepal width

Name	Value
mean	2.77
std	0.31
max	3.40
75%	3.00
50%	2.80
25%	2.53
min	2.00

Table 9: Stats.




Figure 12: Distribution

## Problem state VIII




Figure 13: Petal length

Name	Value
mean	4.26
std	0.47
max	5.10
75%	4.60
50%	4.35
25%	4.00
min	3.00

Table 10: Stats.




Figure 14: Distribution

## Problem state IX




Figure 15: Petal width

Name	Value
mean	1.33
std	0.20
max	1.80
75%	1.50
50%	1.30
25%	1.20
min	1.00

Table 11: Stats.

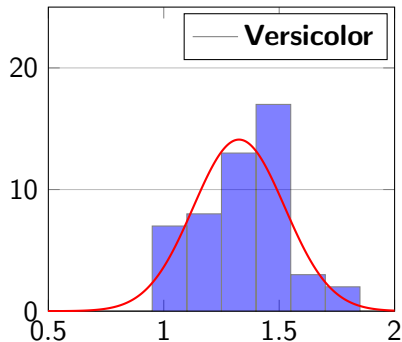


Figure 16: Distribution

## Problem state X




Figure 17: Sepal length

Name	Value
mean	6.59
std	0.64
max	7.90
75%	6.90
50%	6.50
25%	6.23
min	4.90

Table 12: Stats.




Figure 18: Distribution



## Problem state XI




Figure 19: Sepal width

Name	Value
mean	2.97
std	0.32
max	3.80
75%	3.18
50%	3.00
25%	2.80
min	2.20

Table 13: Stats.




Figure 20: Distribution

## Problem state XII




Figure 21: Petal length

Name	Value
mean	5.55
std	0.55
max	6.90
75%	5.88
50%	5.55
25%	5.10
min	4.50

Table 14: Stats.

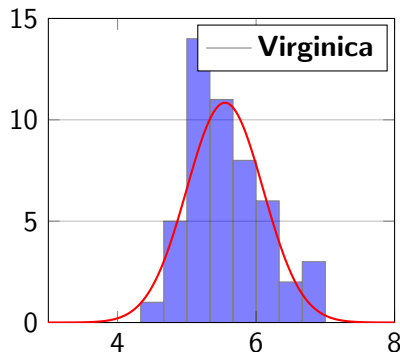


Figure 22: Distribution

## Problem state XIII




Figure 23: Petal width

Name	Value
mean	2.03
std	0.27
max	2.50
75%	2.30
50%	2.00
25%	1.80
min	1.40

Table 15: Stats.



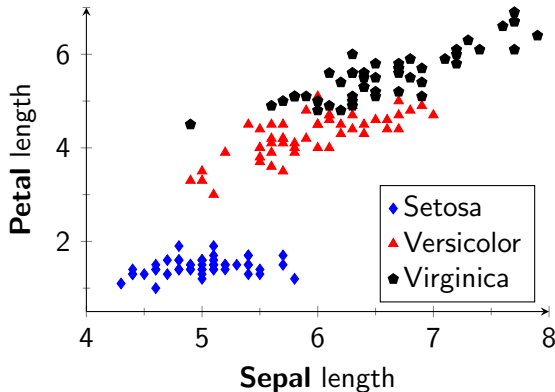


Figure 24: Distribution

## Scatter plot (left) and PCA (right) on iris dataset



# Shapiro – Wilk's normality test ([Razali et al., 2011])

**Prerequisites:**

**H0** normal.

**H1** not normal.

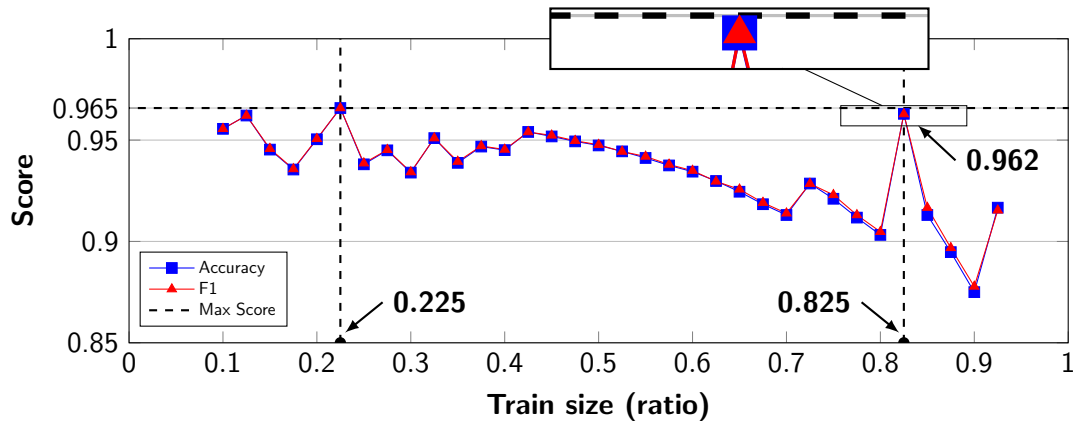
$\alpha$  10%.

	Setosa	Versicolor	Virginica
<b>Sepal len.</b>	0.459	0.465	0.258
<b>Sepal wid.</b>	0.272	0.338	0.181
<b>Petal len.</b>	0.055	0.158	0.110
<b>Petal wid.</b>	0.000	0.027	0.087

Table 16: P-val for SW test statistics

**Conclusion:** under other things being equal data is normal.

## Fitting results



## Summary: topics covered

- ① **Purposes** and **applications** of Bayesian Classifier.
- ② **Maths** behind the algorithm.
- ③ **Limitations**.
- ④ **Example** on iris dataset.
- ⑤ **Best** score (**0.965** accuracy) reached at **0.225** train ratio.

## References

[Bayes and Price, 1763] Bayes, T. and Price, R. (1763).

An essay towards solving a problem in the doctrine of chances.

*Philosophical Transactions (1683-1775)*, 53:370–418.

[Razali et al., 2011] Razali, N. M., Wah, Y. B., et al. (2011).

Power comparisons of shapiro-wilk, kolmogorov-smirnov, lilliefors and anderson-darling tests.

*Journal of statistical modeling and analytics*, 2(1):21–33.

[Wilks, 1938] Wilks, S. S. (1938).

The large-sample distribution of the likelihood ratio for testing composite hypotheses.

*The annals of mathematical statistics*, 9(1):60–62.



## Conclusions

Thank you for attention!

Questions are welcome.